

1921
— 2021
UN SECOLO
DI STORIA
D'AVANTI A NOI



UNIVERSITÀ
CATTOLICA
del Sacro Cuore

Time Series Analysis and Forecasting

MGO962

Lab 1: Data Manipulation in Python

Main Python packages



Numpy library

- introduces objects for multidimensional arrays and matrices, as well as functions that allow to easily perform advanced mathematical and statistical operations on those objects
- provides vectorization of mathematical operations on arrays and matrices which significantly improves the performance
- many other python libraries are built on NumPy

Scipy library

- collection of algorithms for linear algebra, differential equations, numerical integration, optimization, statistics and more
- part of SciPy Stack
- built on NumPy

Pandas library

- adds data structures and tools designed to work with table-like data (similar to Series and Data Frames in R)
- provides tools for data manipulation: reshaping, merging, sorting, slicing, aggregation etc.
- allows handling missing data

Matplotlib library

- python 2D plotting library which produces publication quality figures in a variety of hardcopy formats
- a set of functionalities similar to those of MATLAB
- line plots, scatter plots, barcharts, histograms, pie charts etc.
- relatively low-level; some effort needed to create advanced visualization

Loading Python Libraries

#Import Python Libraries

```
import numpy as np
```

```
import scipy as sp
```

```
import pandas as pd
```

```
import matplotlib as mpl
```

Read CSV Files

```
import pandas as pd
df = pd.read_csv("data/tourism.csv", parse_dates=True)
df = df.drop(columns=["Unnamed: 0"])
df.head()
```

	Quarter	Region	State	Purpose	Trips
## 0	1998-01-01	Adelaide	South Australia	Business	135.077690
## 1	1998-04-01	Adelaide	South Australia	Business	109.987316
## 2	1998-07-01	Adelaide	South Australia	Business	166.034687
## 3	1998-10-01	Adelaide	South Australia	Business	127.160464
## 4	1999-01-01	Adelaide	South Australia	Business	137.448533

Data Frame Data Types

Pandas Type	Native Python Type	Description
object	string	The most general dtype. Will be assigned to your column if column has mixed types (numbers and strings).
int64	int	Numeric characters. 64 refers to the memory allocated to hold this character.
float64	float	Numeric characters with decimals. If a column contains numbers and NaNs(see below), pandas will default to float64, in case your missing value has a decimal.
datetime64, timedelta[ns]	N/A (but see the datetime module in Python's standard library)	Values meant to hold time data. Look into these for time series experiments.

Data Frame Data Types

```
import pandas as pd
df = pd.read_csv("data/tourism.csv", parse_dates=True)
df = df.drop(columns=["Unnamed: 0"])
df.dtypes
```

```
## Quarter      object
## Region       object
## State        object
## Purpose      object
## Trips        float64
## dtype: object
```

Data Frame Data Attributes

<u>df.attribute</u>	description
dtypes	list the types of the columns
columns	list the column names
axes	list the row labels and column names
ndim	number of dimensions
size	number of elements
shape	return a tuple representing the dimensionality
values	numpy representation of the data

Data Frame Data Methods

<u>df.method()</u>	description
head([n]), tail([n])	first/last n rows
describe()	generate descriptive statistics (for numeric columns only)
max(), min()	return max/min values for all numeric columns
mean(), median()	return mean/median values for all numeric columns
std()	standard deviation
sample([n])	returns a random sample of the data frame
dropna()	drop all the records with missing values

Selecting a column in a Data Frame 1

Method 1: Subset the data frame using column name:

```
import pandas as pd
df = pd.read_csv("data/tourism.csv", parse_dates=True)
df['State']
```

```
## 0      South Australia
## 1      South Australia
## 2      South Australia
## 3      South Australia
## 4      South Australia
##      ...
## 23403   South Australia
## 23404   South Australia
## 23405   South Australia
## 23406   South Australia
## 23407   South Australia
## Name: State, Length: 23408, dtype: object
```

Selecting a column in a Data Frame 2

Method 2: Use the column name as an attribute:

```
import pandas as pd
df = pd.read_csv("data/tourism.csv", parse_dates=True)
df.State
```

```
## 0      South Australia
## 1      South Australia
## 2      South Australia
## 3      South Australia
## 4      South Australia
##      ...
## 23403   South Australia
## 23404   South Australia
## 23405   South Australia
## 23406   South Australia
## 23407   South Australia
## Name: State, Length: 23408, dtype: object
```

Data Frames `groupby` method 1

Using “`groupby`” method we can:

- Split the data into groups based on some criteria
- Calculate statistics (or apply a function) to each group
- Similar to `dplyr()` function in R

Data Frames groupby method 2

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df_state = df.groupby(['rank'])
df_state.mean()
```

	phd	service	salary
## rank			
## AssocProf	15.076923	11.307692	91786.230769
## AsstProf	5.052632	2.210526	81362.789474
## Prof	27.065217	21.413043	123624.804348

Data Frames groupby 3

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df_state = df.groupby(['rank'])[['salary']]
df_state.mean()
```

##	salary
## rank	
## AssocProf	91786.230769
## AsstProf	81362.789474
## Prof	123624.804348

Data Frames groupby 4

groupby performance notes:

- no grouping/splitting occurs until it's needed. Creating the groupby object only verifies that you have passed a valid mapping
- by default the group keys are sorted during the groupby operation. You may want to pass `sort=False` for potential speedup:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df_state = df.groupby(['rank'], sort=False)[['salary']]
df_state.mean()
```

```
##                salary
## rank
## Prof          123624.804348
## AssocProf     91786.230769
## AsstProf      81362.789474
```

Data Frames filtering 1

To subset the data we can apply Boolean indexing. This indexing is commonly known as a filter. For example if we want to subset the rows in which the salary value is greater than \$120K:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df_sub = df[ df['salary'] > 120000 ]
df_sub.mean()
```

```
## phd                28.8
## service            24.6
## salary             141722.4
## dtype: float64
```

Data Frames filtering 2

Any Boolean operator can be used to subset the data:

- ">" greater; >= greater or equal;
- "<" less; <= less or equal;
- "==" equal; != not equal;

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df_f = df[ df['sex'] == 'Female' ]
df_f.mean()
```

```
## phd                16.512821
## service            11.564103
## salary             101002.410256
## dtype: float64
```

Data Frames slicing 1

There are a number of ways to subset the Data Frame:

- one or more columns
- one or more rows
- a subset of rows and columns

Rows and columns can be selected by their position or label

Data Frames slicing 2

When selecting one column, it is possible to use single set of brackets, but the resulting object will be a Series (not a DataFrame):

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df['salary']
```

```
## 0      186960
## 1       93000
## 2      110515
## 3      131205
## 4      104800
##      ...
## 73     105450
## 74     104542
## 75     124312
## 76     109954
## 77     109646
```

Data Frames slicing 3

When we need to select more than one column and/or make the output to be a DataFrame, we should use double brackets:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df[['rank', 'salary']]
```

	rank	salary
## 0	Prof	186960
## 1	Prof	93000
## 2	Prof	110515
## 3	Prof	131205
## 4	Prof	104800
##
## 73	Prof	105450
## 74	AssocProf	104542
## 75	Prof	124312
## 76	Prof	109954

Data Frames slicing 4

If we need to select a range of rows, we can specify the range using ":"

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df[10:15]
```

##	rank	discipline	phd	service	sex	salary
## 10	Prof	B	39	33	Male	128250
## 11	Prof	B	23	23	Male	134778
## 12	AsstProf	B	1	0	Male	88000
## 13	Prof	B	35	33	Male	162200
## 14	Prof	B	25	19	Male	153750

Notice that the first row has a position 0, and the last value in the range is omitted: So for 0:10 range the first 10 rows are returned with the positions starting with 0 and ending with 9

Data Frames loc method

If we need to select a range of rows, using their labels we can use method loc:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df.loc[10:20,['rank','sex','salary']]
```

	rank	sex	salary
## 10	Prof	Male	128250
## 11	Prof	Male	134778
## 12	AsstProf	Male	88000
## 13	Prof	Male	162200
## 14	Prof	Male	153750
## 15	Prof	Male	150480
## 16	AsstProf	Male	75044
## 17	AsstProf	Male	92000
## 18	Prof	Male	107300
## 19	Prof	Male	150500
## 20	AsstProf	Male	92000

Data Frames `iloc` method

If we need to select a range of rows and/or columns, using their positions we can use method `iloc`:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df.iloc[10:20,[0, 3, 4, 5]]
```

	rank	service	sex	salary
## 10	Prof	33	Male	128250
## 11	Prof	23	Male	134778
## 12	AsstProf	0	Male	88000
## 13	Prof	33	Male	162200
## 14	Prof	19	Male	153750
## 15	Prof	3	Male	150480
## 16	AsstProf	3	Male	75044
## 17	AsstProf	0	Male	92000
## 18	Prof	7	Male	107300
## 19	Prof	27	Male	150500

Data Frames `iloc` method 2

We can sort the data by a value in the column. By default the sorting will occur in ascending order and a new data frame is return.

```
df.iloc[0]    # First row of a data frame  
df.iloc[5]    #(i+1)th row  
df.iloc[-1]   # Last row
```

```
df.iloc[:, 0] # First column  
df.iloc[:, -1] # Last column
```

```
df.iloc[0:7]          #First 7 rows  
df.iloc[:, 0:2]       #First 2 columns  
df.iloc[1:3, 0:2]     #Second through third rows and first 2 columns  
df.iloc[[0,5], [1,3]] #1st and 6th rows and 2nd and 4th columns
```

Data Frames sorting method

We can sort the data using 2 or more columns:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
# Create a new data frame sorted by the column Salary
df_sorted = df.sort_values( by ='service')
df_sorted.head()
```

##		rank	discipline	phd	service	sex	salary
## 55	AsstProf		A	2	0	Female	72500
## 23	AsstProf		A	2	0	Male	85000
## 43	AsstProf		B	5	0	Female	77000
## 17	AsstProf		B	4	0	Male	92000
## 12	AsstProf		B	1	0	Male	88000

Data Frames sorting method 2

We can sort the data using 2 or more columns:

```
import pandas as pd
df = pd.read_csv("data/salaries.csv")
df_sorted = df.sort_values( by =['service', 'salary'], ascending =
df_sorted.head(10)
```

##	rank	discipline	phd	service	sex	salary
## 52	Prof	A	12	0	Female	105000
## 17	AsstProf	B	4	0	Male	92000
## 12	AsstProf	B	1	0	Male	88000
## 23	AsstProf	A	2	0	Male	85000
## 43	AsstProf	B	5	0	Female	77000
## 55	AsstProf	A	2	0	Female	72500
## 57	AsstProf	A	3	1	Female	72500
## 28	AsstProf	B	7	2	Male	91300
## 42	AsstProf	B	4	2	Female	80225
## 68	AsstProf	A	4	2	Female	77500

Missing Values

Missing values are marked as NaN

```
import pandas as pd
df = pd.read_csv("data/flights.csv")
```

```
## sys:1: DtypeWarning: Columns (7,8) have mixed types. Specify dt
```

```
df[df.isnull().any(axis=1)].head()
```

```
##      YEAR  MONTH  DAY  ...  AIRLINE_DELAY  LATE_AIRCRAFT_DELAY  WE
## 0   2015      1    1  ...              NaN                  NaN
## 1   2015      1    1  ...              NaN                  NaN
## 2   2015      1    1  ...              NaN                  NaN
## 3   2015      1    1  ...              NaN                  NaN
## 4   2015      1    1  ...              NaN                  NaN
##
## [5 rows x 31 columns]
```

Missing Values 2

<u>df.method()</u>	description
dropna()	Drop missing observations
dropna(how='all')	Drop observations where all cells is NA
dropna(axis=1, how='all')	Drop column if all the values are missing
dropna(thresh = 5)	Drop rows that contain less than 5 non-missing values
fillna(0)	Replace missing values with zeros
isnull()	returns True if the value is missing
notnull()	Returns True for non-missing values

Missing Values 3

- When summing the data, missing values will be treated as zero
- If all values are missing, the sum will be equal to NaN
- `cumsum()` and `cumprod()` methods ignore missing values but preserve them in the resulting arrays
- Missing values in `GroupBy` method are excluded (just like in R)
- Many descriptive statistics methods have `skipna` option to control if missing data should be excluded . This value is set to `True` by default (unlike R)

Aggregation Functions in Pandas

Aggregation - computing a summary statistic about each group, i.e.

- compute group sums or means
- compute group sizes/counts

Common aggregation functions:

- min, max
- count, sum, prod
- mean, median, mode, mad
- std, var

Aggregation Functions in Pandas 2

agg() method are useful when multiple statistics are computed per column:

```
import pandas as pd
df = pd.read_csv("data/flights.csv")
df[['DEPARTURE_DELAY', 'ARRIVAL_DELAY']].agg(['min', 'mean', 'max'])
```

	DEPARTURE_DELAY	ARRIVAL_DELAY
## min	-82.000000	-87.000000
## mean	9.370158	4.407057
## max	1988.000000	1971.000000